



# Is your test plan still **true**?

IBM Bob reads your test plan, spec and release checklist, finds the test behind each claim, then breaks the code to see if the test notices.

— THE PROBLEM

# Green badge. No tests.

A while back I shipped a project whose CI badge was green for days while it ran no tests at all. It was waiting on a credential that didn't exist. Nobody looked, because it was green.

Test plans have the same problem. Someone signs one off before every release, and **nobody really checks.**

| Summary, test-plan.xlsx    |      |
|----------------------------|------|
| Cases in plan              | 90   |
| Automated                  | 84   |
| Automated cases passing    | 84   |
| Automated coverage of plan | 100% |
| Release gate               | Met  |

— WHAT IT IS

# A Bob mode and a skill. **Five stages.**

Copy `.bob/` into your repository, switch Bob to the Plumblin mode, and ask it to audit.

01

## Read

The `.xlsx` plan, `.docx` spec and `.pdf` checklist become claims, each with its source cell.

02

## Map

One subagent per test file finds the test behind each case, even with no id in its name.

03

## Break

A subagent that never sees the test makes the claim false and writes a witness. The runner applies it in a throwaway worktree.

04

## Judge

Every spec clause and checklist tick against the code, the changelog and git history.

05

## Report

One HTML report that leads with what is false.

# The plan says 100%. Here's what Bob found.

Turnstile, a small auth service I wrote for this, audited by Bob in the Plumblin mode. I wrote every defect down in an answer key before Bob ran.

84

AUTOMATED CASES

60

WITH A TEST

24

WITH NO TEST

8

TESTS IN NAME ONLY

4 of 24

SPEC CLAUSES FALSE

4 of 10

CHECKLIST TICKS FALSE

3 of 6

SUMMARY LINES FALSE

— BREAK IT. DOES THE TEST NOTICE?

# TC-41: expired tokens are rejected. Delete the check. Still green.

test/tokens.test.mjs

● PASSES

```
test('TC-41 rejects an expired access token', () => {
  const t = issueAccess('ada@...', ['user'], { now });
  const result = verifyAccess(t, { now: now + 3h });
  assert.ok(result); // always true
});
```

Bob's change, written without seeing the test

```
- if (claims.exp <= t - SKEW) return { ...'expired' };
+ if (false)                  return { ...'expired' };
```

Bob's witness

proves the break

```
const result = verifyAccess(t, { now: now + 2h });
return result.valid === false
      && result.reason === 'expired';
real code: true   broken code: false
```

## Test in name only

The expiry check is gone, the witness proves it, and the test still passes, along with the rest of the suite.

— NO ACCUSATION WITHOUT PROOF

# Mutation testing can blame a good test. **So Bob has to prove it.**

- My first run blamed 4 tests out of 12 for changes that didn't really break anything. So now every accusation needs a witness.
- No witness: **UNPROVEN**. Witness wrong on real code: **WITNESS\_INVALID**. Change harmless: **WEAK\_MUTATION**.
- Strip the witnesses out and the same run accuses nobody.

## BOB AGAINST THE ANSWER KEY

|                                   |         |
|-----------------------------------|---------|
| Cases mapped to the same test     | 83 / 84 |
| Spec clauses judged the same      | 24 / 24 |
| Checklist items judged the same   | 10 / 10 |
| Tests in name only the key missed | 2       |

# Bob didn't just build it. **Plumblineline runs inside Bob.**

01

## Custom mode

The Plumblineline mode: role and allowed tools. `.bob/custom_modes.yaml`

02

## Skills

Five stages, the claim schema, the cost rules. `SKILL.md`

03

## Document understanding

Reads the .xlsx and .docx as they are. It couldn't read the .pdf, so Bob wrote a small reader for it.

04

## Subagents

One per test file to map, one per test file to break, with the test withheld.

05

## Agent mode + commands

Runs the mutation runner, and read-only git to judge the spec against history.

06

## Bob as the builder

Planned and wrote the mode, skill, runner, report and tests, and ran every audit.

— BUILT WITH BOB, DURING THE EVENT

# Built with Bob, **eight tasks.**

**8**

BOB TASKS

**29.86**

OF 40 BOBCOINS

**86**

TESTS, ALL GREEN

**0**

ACCUSATIONS WITHOUT PROOF

- **01 to 03** built the mode, the skill and the runner. **04 to 08** ran the audit and tightened up the skill as problems showed up.
- **I reviewed every task.** Bob's first plan aimed mutations at what the test already checks, and its first go at judging called a false tick "partial". Every correction is logged, with who fixed it.
- **Evidence:** [docs/bob/SESSIONS.md](#), exports in [docs/bob/tasks/](#), screenshots in [bob\\_sessions/](#).

— THE FREE HALF, FOR ANY REPOSITORY

# 23 of 100

of the most-starred installable repositories on GitHub owned by organisations fail a check against their own README. No AI, a negative control for every check, and I confirmed every finding by hand. 25 Sep 2026.

- **react/create-react-app** still sends people to documentation at an address that stopped working.
- **kubernetes/kubernetes** links to a case-studies page that is a 404.
- **laravel/laravel** refers to a licence file that is not there.
- **Paste any public repository:** `plumbline-smoky.vercel.app/check`

# \$20 per repository per month.

## FREE

### The deterministic checks

Hosted and in CI, for any repository. No Bob, no marginal cost. The top of the funnel.

## PAID

### The document tier

Test plan, spec and checklist: every claim mapped, every test proved by breaking the code behind it.

## WHO PAYS

### The release owner

Whoever signs the checklist. They're on the hook when it's wrong, and in regulated teams they have to show an auditor the plan was really tested.

## WHOSE BOB

### The team's own seat

Plumblin installs as a mode and a skill. The Bobcoins are the team's; the \$20 is for Plumblin. It sits in the release path, so one repository leads to the rest.

— DON'T TAKE MY WORD FOR IT

# Check it **yourself.**

|                                      |                                                                    |
|--------------------------------------|--------------------------------------------------------------------|
| Bob's audit of Turnstile             | <code>examples/turnstile/.plumbline/</code> and the live report    |
| Bob against the answer key           | <code>node examples/score.mjs</code>                               |
| The answer key itself                | <code>verify-ground-truth.mjs --self-test</code> : 83 of 83 caught |
| Every figure on every surface agrees | <code>node facts/render.mjs --check</code>                         |
| The free-tier number                 | <code>docs/MEASUREMENT.md</code> , with all eighteen corrections   |

## Is your test plan still **true?**

Solo entry. Previously 3rd of 190 at KeeperHub,  
4th of 100+ at Ready, Spec, Ship.