



Is your test plan still **true**?

IBM Bob reads your test plan, spec and release checklist, finds the test behind each claim, then breaks the code to see if the test notices.

— THE PROBLEM

Green badge. No tests.

A while back I shipped a project whose CI badge was green for days while it ran no tests at all. It was waiting on a credential that didn't exist. Nobody looked, because it was green.

Test plans have the same problem. Someone signs one off before every release, and **nobody really checks.**

Summary, test-plan.xlsx	
Cases in plan	90
Automated	84
Automated cases passing	84
Automated coverage of plan	100%
Release gate	Met

— WHAT IT IS

A Bob mode and a skill. **Five stages.**

Copy `.bob/` into your repository, switch Bob to the Plumblin mode, and ask it to audit.

01

Read

The `.xlsx` plan, `.docx` spec and `.pdf` checklist become claims, each with its source cell.

02

Map

One subagent per test file finds the test behind each case, even with no id in its name.

03

Break

A subagent that never sees the test makes the claim false and writes a witness. The runner applies it in a throwaway worktree.

04

Judge

Every spec clause and checklist tick against the code, the changelog and git history.

05

Report

One HTML report that leads with what is false.

— THE WORKED EXAMPLE

The plan says 100%. Here's what Bob found.

Turnstile, a small auth service I wrote for this, audited by Bob in the Plumblin mode. I wrote every defect down in an answer key before Bob ran.

84

AUTOMATED CASES

60

WITH A TEST

24

WITH NO TEST

8

TESTS IN NAME ONLY

4 of 24

SPEC CLAUSES FALSE

4 of 10

CHECKLIST TICKS FALSE

3 of 6

SUMMARY LINES FALSE

— BREAK IT. DOES THE TEST NOTICE?

TC-41: expired tokens are rejected. Delete the check. Still green.

test/tokens.test.mjs

● PASSES

```
test('TC-41 rejects an expired access token', () => {
  const t = issueAccess('ada@...', ['user'], { now });
  const result = verifyAccess(t, { now: now + 3h });
  assert.ok(result); // always true
});
```

Bob's change, written without seeing the test

```
- if (claims.exp <= t - SKEW) return { ...'expired' };
+ if (false)                  return { ...'expired' };
```

Bob's witness

proves the break

```
const result = verifyAccess(t, { now: now + 2h });
return result.valid === false
      && result.reason === 'expired';
real code: true   broken code: false
```

Test in name only

The expiry check is gone, the witness proves it, and the test still passes, along with the rest of the suite.

— NO ACCUSATION WITHOUT PROOF

Mutation testing can blame a good test. **So Bob has to prove it.**

- My first run blamed 4 tests out of 12 for changes that didn't really break anything. So now every accusation needs a witness.
- No witness: **UNPROVEN**. Witness wrong on real code: **WITNESS_INVALID**. Change harmless: **WEAK_MUTATION**.
- Strip the witnesses out and the same run accuses nobody.

BOB AGAINST THE ANSWER KEY

Cases mapped to the same test	83 / 84
Spec clauses judged the same	24 / 24
Checklist items judged the same	10 / 10
Tests in name only the key missed	2

Bob didn't just build it. **Plumblineline runs inside Bob.**

01

Custom mode

The Plumblineline mode: role and allowed tools. `.bob/custom_modes.yaml`

02

Skills

Five stages, the claim schema, the cost rules. `SKILL.md`

03

Document understanding

Reads the .xlsx and .docx as they are. It couldn't read the .pdf, so Bob wrote a small reader for it.

04

Subagents

One per test file to map, one per test file to break, with the test withheld.

05

Agent mode + commands

Runs the mutation runner, and read-only git to judge the spec against history.

06

Bob as the builder

Planned and wrote the mode, skill, runner, report and tests, and ran every audit.

— BUILT WITH BOB, DURING THE EVENT

Built with Bob, **eight tasks.**

8

BOB TASKS

29.86

OF 40 BOBCOINS

86

TESTS, ALL GREEN

0

ACCUSATIONS WITHOUT PROOF

- **01 to 03** built the mode, the skill and the runner. **04 to 08** ran the audit and tightened up the skill as problems showed up.
- **I reviewed every task.** Bob's first plan aimed mutations at what the test already checks, and its first go at judging called a false tick "partial". Every correction is logged, with who fixed it.
- **Evidence:** [docs/bob/SESSIONS.md](#), exports in [docs/bob/tasks/](#), screenshots in [bob_sessions/](#).

— THE FREE HALF, FOR ANY REPOSITORY

23 of 100

of the most-starred installable repositories on GitHub owned by organisations fail a check against their own README. No AI, a negative control for every check, and I confirmed every finding by hand. 25 Sep 2026.

- **react/create-react-app** still sends people to documentation at an address that stopped working.
- **kubernetes/kubernetes** links to a case-studies page that is a 404.
- **laravel/laravel** refers to a licence file that is not there.
- **Paste any public repository:** `plumbline-smoky.vercel.app/check`

— BUSINESS

\$20 per repository per month.

FREE

The deterministic checks

Hosted and in CI, for any repository. No Bob, no marginal cost. The top of the funnel.

PAID

The document tier

Test plan, spec and checklist: every claim mapped, every test proved by breaking the code behind it.

WHO PAYS

The release owner

Whoever signs the checklist. They're on the hook when it's wrong, and in regulated teams they have to show an auditor the plan was really tested.

WHOSE BOB

The team's own seat

Plumblineline installs as a mode and a skill. The Bobcoins are the team's; the \$20 is for Plumblineline. It sits in the release path, so one repository leads to the rest.

— DON'T TAKE MY WORD FOR IT

Check it **yourself.**

Bob's audit of Turnstile	<code>examples/turnstile/.plumbline/</code> and the live report
Bob against the answer key	<code>node examples/score.mjs</code>
The answer key itself	<code>verify-ground-truth.mjs --self-test</code> : 83 of 83 caught
Every figure on every surface agrees	<code>node facts/render.mjs --check</code>
The free-tier number	<code>docs/MEASUREMENT.md</code> , with all eighteen corrections

Is your test plan still **true?**

Solo entry.
IBM Bob 2.0 Hackathon.