

Plumblineline

It reads what your project says about itself, then goes and checks whether any of it is true.

18 of the 100 most-starred installable repositories on GitHub fail at least one check. That is 14,304,526 stars between them.

18 of 100. Every one of them wrong about something.

11 README links that go nowhere
every one re-checked with a real request, not a HEAD

3 manifests whose version is not the version on the registry
counted only where the registry points back at the repo, so a shared name is not a finding

2 links pointing at files that are not in the repository
you click, and get nothing

Measured 17 September 2026 against the strongest end of GitHub, not a random sample. My first pass said 89 of 100. That was wrong in the direction that flattered the tool, so I re-checked every finding by hand. Five faults in my own checks, and one in my own process, are written up in the method. The number came down 89, then 47, then 30, then 22, then 19, then 18. Every move was downward, because every move removed an accusation the evidence did not support.

The hour before you ship.

You think you are finished. Nobody has read the README since you wrote it, CI has been green for a month, and the demo link in it was renamed in March.

Green is not the same as true. Nothing in a normal pipeline reads what the documentation claims and then goes to check. So the claims rot quietly, and the first person to find out is a user.

Green is not the same as true.

Three parts. The third one is the point.

CLAIMS

What the project says

Read out of the prose and the documents. The README, the manifest, the specification, the release checklist, the test plan.

CHECKS

What is actually true

Found out separately from the claim. A registry lookup, a real web request, a file that either exists or does not.

PROOF


That the check can fail

A deliberate change that **must** make the check fail. If it does not, the check is reported as unproven. Never as a pass.

A check nobody has ever seen fail is not evidence. It is a green tick, and a green tick can mean nothing at all.

Plumblne

It checks whether a repository does what it says it does, and proves that each of its own checks can fail.



The half that needs no judgement has no Bob dependency and runs in CI on every push.

One README. Two filenames. One of them never existed.

```
[ADR-079] (docs/adr/ADR-079-camera-ground-truth-training.md)    exists
[ADR-079] (docs/adr/ADR-079-camera-supervised-pose-finetune.md) never existed
```

A reader clicks that and gets a 404. The author never saw it, because they never clicked their own link after renaming the file. No test catches this. Nothing in CI reads the README.

ruvnet/RuView, checked against the GitHub contents API. Not a made-up example.

Every check proves it can fail.

Link checkers exist. Drift detectors exist, and they only see what a commit touched, so a README that was wrong when it was written stays invisible. What none of them do is prove their own checks can fail before they report a pass.

I have shipped a green CI badge that was doing nothing at all. A conformance job ran nightly for days and skipped every test, because it waited on a credential that did not exist. It passed, so nothing drew attention to it. I found it on submission day.

```
$ npx github:iamrobertmoore/plumblin --selfcheck
6/6 checks proved able to fail. Result: PROVEN.
check          good input  broken input  what the broken input said
link.relative   accepted   rejected     missing README.md
link.external   accepted   rejected     ERR plumblin.invalid/__control__
manifest.version accepted   rejected     0.0.0-plumblin-control vs 1.3.0
a check that cannot tell those two apart is reported UNPROVEN, never a pass
```

Everything here that needs judgement is Bob.

DOCUMENT UNDERSTANDING

A claim can live in a .docx, a .pdf or a .xlsx, not just in markdown. A specification, a release checklist, a test plan. Bob reads those as they are.

SUBAGENTS

One subagent per claim, each with its own clean context, so nothing gets checked by accident while something else is being checked.

PARALLEL TASKS

Claims that do not depend on each other are checked at the same time. A hundred repositories do not have to queue up.

AGENT MODE

Runs the whole thing, decides what is left to check, and writes the result.

Bob access starts on 25 September, so this is the design. The half that needs no judgement is built, has **no Bob dependency at all**, and runs in CI on every push. It is what produced the number on slide 2.

Your test plan says 84 cases. Your suite covers 61.

Bob reads the spreadsheet the way it is, traces every case to a test that exists, and names the ones nobody wrote.

No other tool can do this, because no other tool reads the plan. Every other check in Plumblin has something sitting next to it that does roughly the same job. This one has nothing.

THE CLAIM

A test plan is a document that claims a level of coverage.

THE CHECK

Trace every case to a test that exists in the suite.

WHAT YOU GET

The cases nobody wrote, named.

Release and deployment is one of the workflows they asked for.

- **The faults are invisible on your machine and obvious to anyone who opens the repo.** So users find them, not the team.
- **The half that needs no AI runs on every push,** at no cost and with no credentials.
- **The result is one self-contained page** you can hand to someone who was not in the room.
- **Nothing in a normal pipeline does this.** Linters read code. Nothing reads what the documentation claims.

MEASURED

3,311 links

fetches and tests across 100 repositories, in **about five minutes**. 1,667 over the network, 1,644 against the repositories' own file trees.

About 3 seconds a repository. By hand this is a day of work, and it is the day nobody spends, which is exactly why the faults survive.

A failing check is not broken software.

A dead link in a README does not stop anyone installing the package. That is exactly why these faults survive: no test reads what the documentation claims.

The number is a snapshot, which is why it carries its date. The sample is deliberately the strongest end of GitHub, so 18 in 100 is a floor rather than an average. Workflows that can pass without checking anything are reported as warnings and kept out of the headline, because switching a docs job to `continue-on-error` is a choice someone made, not a lie.

If the tool cannot see something, it says so. It never calls it a pass and never calls it a failure. It does not guess, and it does not accuse.

Find out what you are actually shipping.

Live report: iamrobertmoore.github.io/plumblines

Source: github.com/iamrobertmoore/plumblines

Method: [docs/MEASUREMENT.md](#)

```
npx github:iamrobertmoore/plumblines --repo .
```

```
npx github:iamrobertmoore/plumblines --selfcheck
```

Exits 0 on a pass or a warning, 1 on a blocker, 2 when a check cannot be shown to fail. No dependencies. It reads your repository and writes nothing.