



Is your test plan still **true**?

IBM Bob reads the test plan, the spec and the release checklist, finds the test behind every claim, then breaks the code to prove the test is real.

— THE PROBLEM

Green badge. No tests.

I once shipped a project whose CI badge was green for days while it ran no tests at all. The job skipped every test, waiting on a credential that did not exist. It passed, so nobody looked.

A test plan is the same promise, in a spreadsheet. Somebody signs it before every release. **Nothing** checks it.

Summary, test-plan.xlsx	
Cases in plan	90
Automated	84
Automated cases passing	84
Automated coverage of plan	100%
Release gate	Met

— WHAT IT IS

A Bob mode and a skill. **Five stages.**

Copy `.bob/` into a repository, open it in IBM Bob, pick the Plumblin mode and ask it to audit.

01

Read

The `.xlsx` plan, `.docx` spec and `.pdf` checklist become claims, each with its source cell.

02

Map

One subagent per test file finds the test behind each case, even with no id in its name.

03

Break

A subagent that never sees the test makes the claim false and writes a witness. The runner applies it in a throwaway worktree.

04

Judge

Every spec clause and checklist tick against the code, the changelog and git history.

05

Report

One HTML report that leads with what is false.

— THE WORKED EXAMPLE

The plan says 100%. Here is what is true.

Turnstile, an auth service written for this entry, audited by Bob in the Plumblin mode. Its defects were recorded in an answer key before Bob ran.

84

AUTOMATED CASES

60

WITH A TEST

24

WITH NO TEST

8

TESTS IN NAME ONLY

4 of 24

SPEC CLAUSES FALSE

4 of 10

CHECKLIST TICKS FALSE

3 of 6

SUMMARY LINES FALSE

— BREAK IT. DOES THE TEST NOTICE?

TC-41: expired tokens are rejected. Delete the check. Still green.

test/tokens.test.mjs

● PASSES

```
test('TC-41 rejects an expired access token', () => {
  const t = issueAccess('ada@...', ['user'], { now });
  const result = verifyAccess(t, { now: now + 3h });
  assert.ok(result); // always true
});
```

Bob's change, written without seeing the test

```
- if (claims.exp <= t - SKEW) return { ...'expired' };
+ if (false)                  return { ...'expired' };
```

Bob's witness

proves the break

```
const result = verifyAccess(t, { now: now + 2h });
return result.valid === false
      && result.reason === 'expired';
real code: true   broken code: false
```

Test in name only

The behaviour is gone, the witness proves it, and the test still passes. So does the whole suite.

— NO ACCUSATION WITHOUT PROOF

Every mutation tester can be fooled. **This one has to prove it was not.**

- In the first run, 4 of 12 surviving changes did not really break the claim. So every accusation now needs a witness.
- No witness: **UNPROVEN**. Witness wrong on real code: **WITNESS_INVALID**. Change harmless: **WEAK_MUTATION**.
- Remove every witness and the same run accuses nobody.

BOB AGAINST THE ANSWER KEY

Cases mapped to the same test	83 / 84
Spec clauses judged the same	24 / 24
Checklist items judged the same	10 / 10
Tests in name only the key missed	2

Remove any one and the document tier **stops working.**

01

Custom mode

The Plumline mode: role and allowed tools. `.bob/custom_modes.yaml`

02

Skills

Five stages, the claim schema, the cost rules. `SKILL.md`

03

Document understanding

Reads .xlsx and .docx as they are. For the .pdf, Bob wrote its own zero-dependency reader.

04

Subagents

One per test file to map, one per test file to break, with the test withheld.

05

Agent mode + commands

Runs the mutation runner, and read-only git to judge the spec against history.

06

Bob as the builder

Planned and wrote the mode, skill, runner, report and tests, and ran every audit.

— BUILT WITH BOB, DURING THE EVENT

Eight Bob tasks. **Every one exported.**

8

BOB TASKS

30.23

OF 40 BOBCOINS

86

TESTS, ALL GREEN

0

ACCUSATIONS WITHOUT PROOF

- **Plan, build, audit, correct.** Tasks 01 to 03 built the mode, skill and runner; 04 to 08 ran the audit and hardened the skill.
- **Reviewing Bob was part of using it.** Its first plan aimed mutations at what the test already asserts; its first judging called a false tick "partial". Every correction is logged, with who fixed it.
- **Evidence:** [docs/bob/SESSIONS.md](#), exports in [docs/bob/tasks/](#), screenshots in [bob_sessions/](#).

— THE FREE HALF, FOR ANY REPOSITORY

23 of 100

of the most-starred installable repositories on GitHub owned by organisations fail a check against their own README. No AI, every check with a negative control, every candidate confirmed by hand. 25 Sep 2026.

- **react/create-react-app** sends readers to documentation at an address that no longer exists.
- **kubernetes/kubernetes** links to a case-studies page that is a 404.
- **laravel/laravel** refers to a licence file that is not there.
- **Paste any public repository:** `plumbline-smoky.vercel.app/check`

\$20 per repository per month.

FREE

The deterministic checks

Hosted and in CI, for any repository. No Bob, no marginal cost. The top of the funnel.

PAID

The document tier

Test plan, spec and checklist: every claim mapped, every test proved by breaking the code behind it.

WHO PAYS

The release owner

The person who signs the checklist, and is on the hook when it is wrong.

WHOSE BOB

The team's own seat

Plumblineline installs as a mode and a skill. The Bobcoins are the team's; the \$20 is for Plumblineline. It sits in the release path, so one repository leads to the rest.

— CHECK EVERY CLAIM IN FIVE MINUTES

Proved, **not asserted.**

Bob's audit of Turnstile	<code>examples/turnstile/.plumbline/</code> and the live report
Bob against the answer key	<code>node examples/score.mjs</code>
The answer key itself	<code>verify-ground-truth.mjs --self-test</code> : 83 of 83 caught
Every figure on every surface agrees	<code>node facts/render.mjs --check</code>
The free-tier number	<code>docs/MEASUREMENT.md</code> , with all eighteen corrections

Solo entrant. 3rd of 190 at KeeperHub. 4th of 100+ at Ready, Spec, Ship.